

OpenShift Container Platform 4 for Operations: EXERCISES

1. Installation

- a. Adding day-1 kernel arguments
- b. Setting DNS to private
- c. Setting the Ingress Controller to private
- d. Restricting the API server to private
- e. Reviewing the installation log
- f. Getting cluster version, status, and update details
- g. Querying the status of the cluster nodes by using the CLI
- h. Gathering logs from a failed installation
- i. Manually gathering logs with SSH access to your host(s)
- j. Manually gathering logs without SSH access to your host(s)
- k. Getting debug information from the installation program

2. Updating a cluster by using the CLI

3. Post-installation

- a. Post-installation machine configuration tasks
 - i. Checking machine config pool status
 - ii. Configuring chrony time service
 - iii. Adding kernel arguments to nodes
 - iv. Adding a real-time kernel to nodes
 - v. Configuring journald settings
 - vi. Configuring container image registry settings
 - vii. Adding extensions to RHCOS
 - viii. Creating a KubeletConfig CRD to edit kubelet parameters
 - ix. Creating a ContainerRuntimeConfig CR to edit CRI-O parameters
- b. Post-installation cluster tasks
 - i. Scaling a machine set manually
 - ii. Creating default cluster-wide node selectors
 - iii. Moving resources to infrastructure machine sets
 - iv. Creating infrastructure machine sets for production environments
 - v. Sample YAML for a machine set custom resource on AWS
 - vi. Creating an infrastructure node
 - vii. Creating infrastructure machines
 - viii. Deploying the cluster autoscaler
 - ix. Deploying the machine autoscaler
 - x. Enabling etcd encryption
 - xi. Disabling etcd encryption
 - xii. Defragmenting etcd data
 - xiii. Specifying the number of pods that must be up with pod disruption budgets

OpenShift Container Platform 4 for Operations: EXERCISES

- c. Post-installation node tasks
 - i. Creating a MachineHealthCheck resource
 - ii. Setting up CPU Manager
 - iii. Configuring huge pages
 - iv. Adding taints and tolerations
 - v. Adding taints and tolerations using a machine set
 - vi. Binding a user to a node using taints and tolerations
 - vii. Controlling nodes with special hardware using taints and tolerations
 - viii. Removing taints and tolerations
 - ix. Disabling or enforcing CPU limits using CPU CFS quotas
 - x. Disabling overcommitment for a project
 - xi. Configuring garbage collection for containers and images
 - xii. Configuring the maximum number of pods per node
- d. Post-installation network configuration
 - i. Creating a network policy
 - ii. Deleting a network policy
 - iii. Viewing network policies
 - iv. Configuring multitenant isolation by using network policy
- e. Preparing for users
 - i. Viewing cluster roles and bindings
 - ii. Viewing local roles and bindings
 - iii. Adding roles to users
 - iv. Creating a local role
 - v. Creating a cluster role
 - vi. Creating a cluster admin
 - vii. Removing the kubeadmin user

4. Troubleshooting

- a. Gathering data about your cluster
 - i. Gathering data about your cluster for Red Hat Support
 - ii. Gathering data about specific features
 - iii. Obtaining your cluster ID
 - iv. Generating a sosreport archive for an OpenShift Container Platform cluster node
 - v. Querying bootstrap node journal logs
 - vi. Querying cluster node journal logs
 - vii. Collecting a network trace from an OpenShift Container Platform node or container
 - viii. Summarizing cluster specifications through clusterversion

OpenShift Container Platform 4 for Operations: EXERCISES

- b. Troubleshooting installations
 - i. Specifying OpenShift Container Platform installer log levels
 - ii. Troubleshooting openshift-install command issues
 - iii. Monitoring installation progress
 - iv. Gathering bootstrap node diagnostic data
 - v. Investigating master node installation issues
 - vi. Investigating etcd installation issues
 - vii. Investigating master node kubelet and API server issues
 - viii. Investigating worker node installation issues
 - ix. Querying Operator status after installation
 - x. Gathering logs from a failed installation
 - c. Verifying node health
 - i. Reviewing node status, resource usage, and configuration
 - ii. Querying the kubelet's status on a node
 - d. Troubleshooting CRI-O container runtime issues
 - i. Verifying CRI-O runtime engine status
 - ii. Gathering CRI-O journald unit logs
 - e. Troubleshooting Operator issues
 - i. Viewing Operator subscription status using the CLI
 - ii. Querying Operator pod status
 - iii. Gathering Operator logs
 - iv. Disabling Machine Config Operator from automatically rebooting
 - f. Investigating pod issues
 - i. Reviewing pod status
 - ii. Inspecting pod and container logs
 - iii. Accessing running pods
 - iv. Starting debug pods with root access
 - v. Copying files to and from pods and containers
 - g. Troubleshooting storage issues
 - i. Resolving multi-attach errors
 - h. Diagnosing OpenShift CLI (oc) issues
 - i. Specifying OpenShift CLI (oc) log levels
5. Web console
- a. Configuring the web console
 - b. Disabling the web console

OpenShift Container Platform 4 for Operations: EXERCISES

6. Security

- a. Container security
 - i. Watching cluster events
- b. Configuring certificates
 - i. Replacing the default ingress certificate
 - ii. Add an API server named certificate
- c. Viewing the audit log
- d. Configuring the audit log policy

7. Encrypting etcd

- a. Enabling etcd encryption
- b. Disabling etcd encryption

8. Authentication

- a. Configuring identity providers
 - i. Configuring an HTPasswd identity provider
 - 1. Creating an HTPasswd file using Linux
 - 2. Creating the HTPasswd Secret
 - 3. Adding an identity provider to your clusters
 - 4. Updating users for an HTPasswd identity provider
- b. Using RBAC to define and apply permissions
 - i. Viewing cluster roles and bindings
 - ii. Viewing local roles and bindings
 - iii. Adding roles to users
 - iv. Creating a local role
 - v. Creating a cluster role
 - vi. Creating a cluster admin
- c. Removing the kubeadmin user
- d. Understanding and creating service accounts
 - i. Creating service accounts
 - ii. Examples of granting roles to service accounts
 - iii. Using a service account's credentials externally
- e. Managing security context constraints
 - i. Creating security context constraints
 - ii. Role-based access to security context constraints
 - iii. Security context constraints reference commands

OpenShift Container Platform 4 for Operations: EXERCISES

- f. Creating and using config maps
 - i. Creating a config map
 - ii. Creating a config map from a directory
 - iii. Creating a ConfigMap from a file
 - iv. Creating a config map from literal values
 - v. Populating environment variables in containers by using config maps
 - vi. Setting command-line arguments for container commands with ConfigMaps
 - vii. Injecting content into a volume by using config maps

9. Networking

- a. Accessing hosts on Amazon Web Services in an installer-provisioned infrastructure cluster
- b. Expanding the node port range
- c. Networking policy
 - i. Creating a network policy
 - ii. Viewing network policies
 - iii. Editing a network policy
 - iv. Deleting a network policy
 - v. Configuring multitenant isolation by using network policy

10. Storage

- a. Persistent storage using hostPath
 - i. Statically provisioning hostPath volumes
 - ii. Mounting the hostPath share in a privileged pod

11. Operators

- a. Administrator tasks
 - i. Installing from OperatorHub using the CLI
 - ii. Deleting Operators from a cluster using the CLI
 - iii. Viewing Operator subscription status using the CLI

OpenShift Container Platform 4 for Operations: EXERCISES

12. Applications

- a. Quotas
 - i. Creating a quota
 - ii. Creating object count quotas
 - iii. Setting resource quota for extended resources
 - iv. Viewing a quota
 - v. Configuring explicit resource quotas
 - vi. Selecting multiple projects during quota creation
 - vii. Viewing applicable cluster resource quotas
- b. Pruning objects to reclaim resources
 - i. Basic pruning operations
 - ii. Pruning groups
 - iii. Pruning DeploymentConfig objects

13. Machine management

- a. Creating a machine set
- b. Scaling a machine set manually
- c. Modifying a machine set
- d. Deleting a specific machine
- e. Deploying the cluster autoscaler
- f. Deploying the machine autoscaler
- g. Sample YAML for a machine set custom resource on AWS
- h. Binding infrastructure node workloads using taints and tolerations
- i. Moving resources to infrastructure machine sets
- j. Creating a MachineHealthCheck resource

14. Working with nodes

- a. Creating daemonsets
- b. Creating jobs
- c. Creating cron jobs
- d. About listing all the nodes in a cluster
- e. Listing pods on a node in your cluster
- f. Viewing memory and CPU usage statistics on your nodes
- g. Configuring master nodes as schedulable
- h. Deleting nodes from a cluster
- i. Adding kernel arguments to nodes
- j. Modifying Nodes
- k. Configuring the maximum number of pods per node
- l. Rebooting a node using pod anti-affinity
- m. Configuring garbage collection for containers and images

OpenShift Container Platform 4 for Operations: EXERCISES

15. Working with clusters

- a. Viewing events using the CLI
- b. Running the cluster capacity tool on the command line
- c. Running the cluster capacity tool as a job inside a pod
- d. Creating a Limit Range
- e. Viewing a limit
- f. Deleting a Limit Range

16. Backup and restore

- a. Backing up etcd data
- b. Identifying an unhealthy etcd member
- c. Determining the state of the unhealthy etcd member
- d. Replacing an unhealthy etcd member whose machine is not running or whose node is not ready
- e. Replacing an unhealthy etcd member whose etcd pod is crashlooping
- f. Shutting down the cluster
- g. Restarting the cluster
- h. Restoring to a previous cluster state
- i. Recovering from expired control plane certificates

17. CLI tools

- a. Installing the CLI by downloading the binary
- b. Logging in to the CLI
- c. Using the CLI
- d. Getting help
- e. Logging out of the CLI
- f. Basic CLI commands
- g. Application management CLI commands
- h. Troubleshooting and debugging CLI commands
- i. Advanced developer CLI commands
- j. Settings CLI commands
- k. Cluster management CLI commands
- l. Node management CLI commands
- m. Security and policy CLI commands
- n. Maintenance CLI commands
- o. Configuration CLI commands
- p. Usage of oc and kubectl commands